

# Clustering

May 22, 2022

## 0.1 Segmenting customers into clusters -

### 0.1.1 Performing Customer Segmentation on the transactional data to build an efficient marketing model.

```
[1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline

import warnings
warnings.filterwarnings('ignore')
```

```
[2]: df = pd.read_excel('Segmenting_customers_into_clusters_dataset.xlsx')
```

```
[3]: df.head()
```

```
[3]: InvoiceNo StockCode Description Quantity \
0 536365 85123A WHITE HANGING HEART T-LIGHT HOLDER 6
1 536365 71053 WHITE METAL LANTERN 6
2 536365 84406B CREAM CUPID HEARTS COAT HANGER 8
3 536365 84029G KNITTED UNION FLAG HOT WATER BOTTLE 6
4 536365 84029E RED WOOLLY HOTTIE WHITE HEART. 6
```

```
InvoiceDate UnitPrice CustomerID Country
0 2010-12-01 08:26:00 2.55 17850.0 United Kingdom
1 2010-12-01 08:26:00 3.39 17850.0 United Kingdom
2 2010-12-01 08:26:00 2.75 17850.0 United Kingdom
3 2010-12-01 08:26:00 3.39 17850.0 United Kingdom
4 2010-12-01 08:26:00 3.39 17850.0 United Kingdom
```

```
[4]: df.tail()
```

```
[4]: InvoiceNo StockCode Description Quantity \
541904 581587 22613 PACK OF 20 SPACEBOY NAPKINS 12
541905 581587 22899 CHILDREN'S APRON DOLLY GIRL 6
541906 581587 23254 CHILDRENS CUTLERY DOLLY GIRL 4
541907 581587 23255 CHILDRENS CUTLERY CIRCUS PARADE 4
```

541908      581587      22138      BAKING SET 9 PIECE RETROSPOT      3

	InvoiceDate	UnitPrice	CustomerID	Country
541904	2011-12-09 12:50:00	0.85	12680.0	France
541905	2011-12-09 12:50:00	2.10	12680.0	France
541906	2011-12-09 12:50:00	4.15	12680.0	France
541907	2011-12-09 12:50:00	4.15	12680.0	France
541908	2011-12-09 12:50:00	4.95	12680.0	France

### 0.1.2 Dataset Variables are as follows:-

**Invoice No:** Invoice number, a 6-digit integral number uniquely assigned to each transaction. If this code starts with the letter 'c', it indicates a cancellation.

**StockCode:** Product / item code, a 5-digit integral number uniquely assigned to each distinct product.

**Description:** Product / item name.

**Quantity:** The quantities of each product / item per transaction.

**Invoice Date:** Invoice Date and time, the day and time when each transaction was generated.

**UnitPrice:** Unit price, Product price per unit in sterling.

**CustomerID:** Customer number, a 5-digit integral number uniquely assigned to each customer.

**Country:** Country name, the name of the country where each customer resides.

```
[5]: # Exploring our dataset
df.shape
```

```
[5]: (541909, 8)
```

```
[6]: df.dtypes
```

```
[6]: InvoiceNo          object
StockCode          object
Description         object
Quantity           int64
InvoiceDate        datetime64[ns]
UnitPrice          float64
CustomerID         float64
Country            object
dtype: object
```


```
[7]: df.describe()
```

```
[7]:
```

	Quantity	UnitPrice	CustomerID
count	541909.000000	541909.000000	406829.000000
mean	9.552250	4.611114	15287.690570
std	218.081158	96.759853	1713.600303
min	-80995.000000	-11062.060000	12346.000000
25%	1.000000	1.250000	13953.000000
50%	3.000000	2.080000	15152.000000
75%	10.000000	4.130000	16791.000000
max	80995.000000	38970.000000	18287.000000


```
[8]: sns.kdeplot(df['UnitPrice'])
```

```
[8]: <AxesSubplot:xlabel='UnitPrice', ylabel='Density'>
```



```
[9]: sns.kdeplot(df['Quantity'])
```

```
[9]: <AxesSubplot:xlabel='Quantity', ylabel='Density'>
```




In our dataset, we have outliers present in the data as there are negative values observed from the plot.

```
[10]: def abs_val(value, data):
        for i in value:
            quant1 = data[i].quantile(0.25)
            quant2 = data[i].quantile(0.75)
            IQR = quant2 - quant1
            val_max = quant2 + 1.5*IQR
            data = data.loc[(data[i] > 0) & (data[i] < val_max)]
        return data
```

```
[11]: item = ['Quantity', 'UnitPrice']
        frame = abs_val(item, df)
```

```
[12]: def clean(value, data):
        plt.figure(figsize = (12,6), dpi=150)
        for j,k in enumerate(value):
            plt.subplot(1, len(value), j+1)
            sns.kdeplot(data[k])
```

```
[13]: clean(item, frame)
```



```
[14]: frame.describe()
```

```
[14]:
```

	Quantity	UnitPrice	CustomerID
count	437169.000000	437169.000000	321030.000000
mean	4.932159	2.712419	15352.180587
std	4.507273	1.944784	1703.612290
min	1.000000	0.001000	12347.000000
25%	1.000000	1.250000	14049.000000
50%	3.000000	2.080000	15298.000000
75%	8.000000	3.750000	16873.000000
max	23.000000	8.490000	18287.000000

```
[15]: import datetime as dt
frame['Trans_year'] = df['InvoiceDate'].dt.year.astype('category')
frame['Trans_month'] = df['InvoiceDate'].dt.month.astype('category')
frame['Trans_day'] = df['InvoiceDate'].dt.day.astype('category')
frame['Trans_hour'] = df['InvoiceDate'].dt.hour.astype('category')
df['CustomerID'] = df['CustomerID'].astype('category')
```

```
[16]: df = frame
```

```
[17]: df.head()
```

```
[17]:
```

	InvoiceNo	StockCode	Description	Quantity	\
0	536365	85123A	WHITE HANGING HEART T-LIGHT HOLDER	6	
1	536365	71053	WHITE METAL LANTERN	6	
2	536365	84406B	CREAM CUPID HEARTS COAT HANGER	8	
3	536365	84029G	KNITTED UNION FLAG HOT WATER BOTTLE	6	

4	536365	84029E	RED WOOLLY HOTTIE WHITE HEART.	6
---	--------	--------	--------------------------------	---

	InvoiceDate	UnitPrice	CustomerID	Country	Trans_year	\
0	2010-12-01 08:26:00	2.55	17850.0	United Kingdom	2010	
1	2010-12-01 08:26:00	3.39	17850.0	United Kingdom	2010	
2	2010-12-01 08:26:00	2.75	17850.0	United Kingdom	2010	
3	2010-12-01 08:26:00	3.39	17850.0	United Kingdom	2010	
4	2010-12-01 08:26:00	3.39	17850.0	United Kingdom	2010	

	Trans_month	Trans_day	Trans_hour
0	12	1	8
1	12	1	8
2	12	1	8
3	12	1	8
4	12	1	8

```
[18]: df.isnull().sum()
```

```
[18]: InvoiceNo      0
      StockCode    0
      Description  0
      Quantity     0
      InvoiceDate   0
      UnitPrice    0
      CustomerID   116139
      Country      0
      Trans_year   0
      Trans_month   0
      Trans_day     0
      Trans_hour    0
      dtype: int64
```

```
[19]: df['CustomerID'].fillna(df['CustomerID'].mode()[0], inplace = True)
```

```
[20]: df.isnull().sum()
```

```
[20]: InvoiceNo      0
      StockCode    0
      Description  0
      Quantity     0
      InvoiceDate   0
      UnitPrice    0
      CustomerID   0
      Country      0
      Trans_year   0
      Trans_month   0
      Trans_day     0
```

```
Trans_hour      0
dtype: int64
```

We have replaced the missing values of Customer ID by the mode of the feature. Now we will Encode the categorical features of the dataset.

```
[24]: def encode(value):
      for a in value:
          mapped = {}
          obj = list(df[a].unique())
          for m,n in enumerate(obj):
              mapped.update({n:m+1})
          df[a] = df[a].map(mapped)
          df[a] = df[a].astype('int64')
```

```
[25]: items = ['InvoiceNo', 'StockCode', 'Description', 'CustomerID', 'Country']
```

```
[26]: encode(items)
```

```
[27]: df.head()
```

```
[27]:
```

	InvoiceNo	StockCode	Description	Quantity	InvoiceDate	UnitPrice	\
0	1	1	1	6	2010-12-01 08:26:00	2.55	
1	1	2	2	6	2010-12-01 08:26:00	3.39	
2	1	3	3	8	2010-12-01 08:26:00	2.75	
3	1	4	4	6	2010-12-01 08:26:00	3.39	
4	1	5	5	6	2010-12-01 08:26:00	3.39	

	CustomerID	Country	Trans_year	Trans_month	Trans_day	Trans_hour
0	1	1	2010	12	1	8
1	1	1	2010	12	1	8
2	1	1	2010	12	1	8
3	1	1	2010	12	1	8
4	1	1	2010	12	1	8

```
[28]: df = df.drop(['InvoiceDate'], axis = 1)
```

```
[29]: # Feature Scaling

from sklearn.preprocessing import StandardScaler
scale = StandardScaler()
```

```
[30]: scale_data = scale.fit_transform(df)
```

```
[32]: scale_data = pd.DataFrame(scale_data, columns = df.columns)
      scale_data.head()
```

```
[32]: InvoiceNo  StockCode  Description  Quantity  UnitPrice  CustomerID  \
0      0.0    -1.139680    -1.152371  0.236915   -0.083515   -0.905703
1      0.0    -1.138779    -1.151482  0.236915    0.348410   -0.905703
2      0.0    -1.137878    -1.150593  0.680643    0.019324   -0.905703
3      0.0    -1.136977    -1.149704  0.236915    0.348410   -0.905703
4      0.0    -1.136076    -1.148815  0.236915    0.348410   -0.905703

      Country  Trans_year  Trans_month  Trans_day  Trans_hour
0 -0.190424    -3.42424    1.254859   -1.616542    -2.13674
1 -0.190424    -3.42424    1.254859   -1.616542    -2.13674
2 -0.190424    -3.42424    1.254859   -1.616542    -2.13674
3 -0.190424    -3.42424    1.254859   -1.616542    -2.13674
4 -0.190424    -3.42424    1.254859   -1.616542    -2.13674
```

### 0.1.3 Segmenting the customers into Clusters with K-Means


```
[33]: from sklearn.cluster import KMeans
```

```
[35]: distance = []
      for i in range(1,20):
          kmeans = KMeans(n_clusters = i)
          kmeans.fit(scale_data)
          distance.append(kmeans.inertia_)
```

```
[36]: plt.figure(figsize=(12,10),dpi=150)
      plt.plot(range(1,20),distance, marker='o')
      plt.xlabel('No. of Clusters')
      plt.ylabel('Inertia')
      plt.title('Elbow Curve')
```

```
[36]: Text(0.5, 1.0, 'Elbow Curve')
```





From the Elbow Curve, it can be observed that after cluster 7 inertia of the remaining clusters is almost constant and hence the number of clusters present in the data can be chosen as 7.

```
[37]: kmeans = KMeans(n_clusters = 7)
      kmeans.fit(scale_data)
      k_pred = kmeans.predict(scale_data)
```

```
[38]: k_pred
```

```
[38]: array([4, 4, 4, ..., 5, 5, 5])
```

```
[40]: scale_data['Cluster_no'] = k_pred
```

```
[41]: scale_data.head()
```

```
[41]:   InvoiceNo  StockCode  Description  Quantity  UnitPrice  CustomerID  \
0         0.0   -1.139680   -1.152371   0.236915   -0.083515   -0.905703
1         0.0   -1.138779   -1.151482   0.236915    0.348410   -0.905703
2         0.0   -1.137878   -1.150593   0.680643    0.019324   -0.905703
```

3	0.0	-1.136977	-1.149704	0.236915	0.348410	-0.905703
4	0.0	-1.136076	-1.148815	0.236915	0.348410	-0.905703

	Country	Trans_year	Trans_month	Trans_day	Trans_hour	Cluster_no
0	-0.190424	-3.42424	1.254859	-1.616542	-2.13674	4
1	-0.190424	-3.42424	1.254859	-1.616542	-2.13674	4
2	-0.190424	-3.42424	1.254859	-1.616542	-2.13674	4
3	-0.190424	-3.42424	1.254859	-1.616542	-2.13674	4
4	-0.190424	-3.42424	1.254859	-1.616542	-2.13674	4

```
[42]: scale_data['Cluster_no'].value_counts()
```


```
[42]: 6    120488
      5     95618
      3     68353
      2     55950
      1     55147
      4     34146
      0       7467
      Name: Cluster_no, dtype: int64
```

```
[56]: print(kmeans.cluster_centers_)
```

```
[[ 0.00000000e+00  2.65964335e-02  2.78644355e-02  6.59217475e-01
   5.54969996e-03  5.48335526e-01  6.63494875e+00  1.88501571e-01
  -1.81997186e-01  1.44936872e-02 -4.32035307e-01]
 [ 0.00000000e+00 -5.34143123e-01 -5.05233035e-01 -1.01289474e-01
  -2.65761983e-01  1.63961212e+00 -1.21526574e-01  2.92035628e-01
   5.95841672e-01  3.97701344e-02  2.98482156e-02]
 [ 0.00000000e+00 -4.14371556e-01 -4.10096508e-01 -5.38554748e-01
   1.62033421e+00 -3.86226776e-01 -1.22897475e-01  2.92035628e-01
  -4.67761206e-01  8.47083275e-02 -2.46670688e-01]
 [ 0.00000000e+00 -4.87913074e-01 -4.63925369e-01  1.44361665e+00
  -5.64948691e-01 -1.19483629e-02 -1.80348918e-02  2.92035628e-01
  -4.65199320e-01  5.33701444e-02 -4.99917614e-01]
 [ 0.00000000e+00 -4.98768369e-01 -5.25752989e-01 -1.70114484e-01
   1.67918014e-01 -7.04513983e-01 -1.33859744e-01 -3.42423973e+00
   1.25485859e+00 -5.75130670e-01  1.00602947e-02]
 [ 0.00000000e+00  1.58374204e+00  1.55190197e+00  5.24522697e-02
  -9.67132036e-03  2.19405460e-01 -9.96604905e-02  2.92035628e-01
   3.42216599e-01  5.21963823e-02 -3.13213686e-02]
 [ 0.00000000e+00 -4.03308787e-01 -3.99309048e-01 -5.57041243e-01
  -3.51688895e-01 -5.73631503e-01 -1.71234909e-01  2.92035628e-01
  -4.07799491e-01  3.28218631e-02  4.33830560e-01]]
```

```
[66]: plt.figure(figsize = (12,10), dpi=150)
      sns.countplot(scale_data['Cluster_no'])
      plt.title('Count plot of Clusters')
```

```
[66]: Text(0.5, 1.0, 'Count plot of Clusters')
```



As we can see the above plot describes the number of datapoints present in each of the clusters.

```
[ ]:
```